

Exercise_3

Hai Nguyen

2023-03-28

Load data

Load the following data: + applications from app_data_sample.parquet + edges from edges_sample.csv

```
applications <- read_parquet("C:/Users/nguye/OneDrive/Documents/MMA/Winter
2023/Network
analysis/672_project_data/app_data_sample.parquet", as_data_frame=TRUE)
edges <- read_csv("C:/Users/nguye/OneDrive/Documents/MMA/Winter 2023/Network
analysis/672_project_data/edges_sample.csv")
```

```
## Rows: 32906 Columns: 4
## — Column specification
```

```
## Delimiter: ","
## chr (1): application_number
## dbl (2): ego_examiner_id, alter_examiner_id
## date (1): advice_date
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.
```

```
applications
```

```
## # A tibble: 2,018,477 × 16
##   applicat...1 filing_d...2 exami...3 exami...4 exami...5 exami...6 exami...7 uspc_...8
uspc_...9
##   <chr>      <date>      <chr>    <chr>    <chr>      <dbl>    <dbl> <chr>
<chr>
## 1 08284457 2000-01-26 HOWARD  JACQUE... V          96082    1764 508
273000
## 2 08413193 2000-10-11 YILDIR... BEKIR    L          87678    1764 208
179000
## 3 08531853 2000-05-17 HAMILT... CYNTHIA <NA>      63213    1752 430
271100
## 4 08637752 2001-07-20 MOSHER  MARY     <NA>      73788    1648 530
388300
## 5 08682726 2000-04-10 BARR    MICHAEL E          77294    1762 427
430100
## 6 08687412 2000-04-28 GRAY    LINDA    LAMEY     68606    1734 156
204000
```

```
## 7 08716371 2004-01-26 MCMILL... KARA RENITA 89557 1627 424
401000
## 8 08765941 2000-06-23 FORD VANESSA L 97543 1645 424
001210
## 9 08776818 2000-02-04 STRZEL... TERESA E 98714 1637 435
006000
## 10 08809677 2002-02-20 KIM SUN U 65530 1723 210
645000
## # ... with 2,018,467 more rows, 7 more variables: patent_number <chr>,
## # patent_issue_date <date>, abandon_date <date>, disposal_type <chr>,
## # appl_status_code <dbl>, appl_status_date <chr>, tc <dbl>, and
abbreviated
## # variable names ^application_number, ^filing_date, ^examiner_name_last,
## # ^examiner_name_first, ^examiner_name_middle, ^examiner_id,
## # ^examiner_art_unit, ^uspc_class, ^uspc_subclass

edges

## # A tibble: 32,906 × 4
## application_number advice_date ego_examiner_id alter_examiner_id
## <chr> <date> <dbl> <dbl>
## 1 09402488 2008-11-17 84356 66266
## 2 09402488 2008-11-17 84356 63519
## 3 09402488 2008-11-17 84356 98531
## 4 09445135 2008-08-21 92953 71313
## 5 09445135 2008-08-21 92953 93865
## 6 09445135 2008-08-21 92953 91818
## 7 09479304 2008-12-15 61767 69277
## 8 09479304 2008-12-15 61767 92446
## 9 09479304 2008-12-15 61767 66805
## 10 09479304 2008-12-15 61767 70919
## # ... with 32,896 more rows
```

Get gender for examiners

We'll get gender based on the first name of the examiner, which is recorded in the field `examiner_name_first`. We'll use library `gender` for that, relying on a modified version of their own [example](#).

Note that there are over 2 million records in the applications table – that's because there are many records for each examiner, as many as the number of applications that examiner worked on during this time frame. Our first step therefore is to get all *unique* names in a separate list `examiner_names`. We will then guess gender for each one and will join this table back to the original dataset. So, let's get names without repetition:

```
library(gender)
#install_genderdata_package() # only run this line the first time you use the
package, to get data for it
# get a list of first names without repetitions
examiner_names <- applications %>%
```

```
distinct(examiner_name_first)
examiner_names
```

```
## # A tibble: 2,595 × 1
##   examiner_name_first
##   <chr>
## 1 JACQUELINE
## 2 BEKIR
## 3 CYNTHIA
## 4 MARY
## 5 MICHAEL
## 6 LINDA
## 7 KARA
## 8 VANESSA
## 9 TERESA
## 10 SUN
## # ... with 2,585 more rows
```

Now let's use function `gender()` as shown in the example for the package to attach a gender and probability to each name and put the results into the table `examiner_names_gender`

```
# get a table of names and gender
examiner_names_gender <- examiner_names %>%
  do(results = gender(.$examiner_name_first, method = "ssa")) %>%
  unnest(cols = c(results), keep_empty = TRUE) %>%
  select(
    examiner_name_first = name,
    gender,
    proportion_female
  )
examiner_names_gender
```

```
## # A tibble: 1,822 × 3
##   examiner_name_first gender proportion_female
##   <chr>               <chr>             <dbl>
## 1 AARON               male             0.0082
## 2 ABDEL               male             0
## 3 ABDOU               male             0
## 4 ABDUL               male             0
## 5 ABDULHAKIM          male             0
## 6 ABDULLAH            male             0
## 7 ABDULLAHI           male             0
## 8 ABIGAIL             female           0.998
## 9 ABIMBOLA            female           0.944
## 10 ABRAHAM             male             0.0031
## # ... with 1,812 more rows
```

Finally, let's join that table back to our original applications data and discard the temporary tables we have just created to reduce clutter in our environment.

```

# remove extra columns from the gender table
examiner_names_gender <- examiner_names_gender %>%
  select(examiner_name_first, gender)
# joining gender back to the dataset
applications <- applications %>%
  left_join(examiner_names_gender, by = "examiner_name_first")
# cleaning up
rm(examiner_names)
rm(examiner_names_gender)
gc()

##           used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  4671496 249.5   8229292 439.5  5114501 273.2
## Vcells 49878864 380.6   93055462 710.0  80194619 611.9

```

Guess the examiner's race

We'll now use package `wru` to estimate likely race of an examiner. Just like with gender, we'll get a list of unique names first, only now we are using surnames.

```

library(wru)
examiner_surnames <- applications %>%
  select(surname = examiner_name_last) %>%
  distinct()
examiner_surnames

## # A tibble: 3,806 × 1
##   surname
##   <chr>
## 1 HOWARD
## 2 YILDIRIM
## 3 HAMILTON
## 4 MOSHER
## 5 BARR
## 6 GRAY
## 7 MCMILLIAN
## 8 FORD
## 9 STRZELECKA
## 10 KIM
## # ... with 3,796 more rows

```

We'll follow the instructions for the package outlined here

<https://github.com/kosukeimai/wru>.

```

examiner_race <- predict_race(voter.file = examiner_surnames, surname.only =
T) %>%
  as_tibble()

## Warning: Unknown or uninitialised column: `state`.

## Proceeding with last name predictions...

```

```
## i All local files already up-to-date!

## 701 (18.4%) individuals' last names were not matched.

examiner_race

## # A tibble: 3,806 × 6
##   surname    pred.whi pred.bla pred.his pred.asi pred.oth
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 HOWARD      0.597    0.295    0.0275   0.00690   0.0741
## 2 YILDIRIM    0.807    0.0273   0.0694   0.0165    0.0798
## 3 HAMILTON    0.656    0.239    0.0286   0.00750   0.0692
## 4 MOSHER      0.915    0.00425  0.0291   0.00917   0.0427
## 5 BARR        0.784    0.120    0.0268   0.00830   0.0615
## 6 GRAY        0.640    0.252    0.0281   0.00748   0.0724
## 7 MCMILLIAN   0.322    0.554    0.0212   0.00340   0.0995
## 8 FORD        0.576    0.320    0.0275   0.00621   0.0697
## 9 STRZELECKA 0.472    0.171    0.220    0.0825    0.0543
## 10 KIM        0.0169   0.00282  0.00546  0.943     0.0319
## # ... with 3,796 more rows
```

As you can see, we get probabilities across five broad US Census categories: white, black, Hispanic, Asian and other. (Some of you may correctly point out that Hispanic is not a race category in the US Census, but these are the limitations of this package.)

Our final step here is to pick the race category that has the highest probability for each last name and then join the table back to the main applications table. See this example for comparing values across columns: <https://www.tidyverse.org/blog/2020/04/dplyr-1-0-0-rowwise/>. And this one for case_when() function: https://dplyr.tidyverse.org/reference/case_when.html.

```
examiner_race <- examiner_race %>%
  mutate(max_race_p = pmax(pred.asi, pred.bla, pred.his, pred.oth, pred.whi))
%>%
  mutate(race = case_when(
    max_race_p == pred.asi ~ "Asian",
    max_race_p == pred.bla ~ "black",
    max_race_p == pred.his ~ "Hispanic",
    max_race_p == pred.oth ~ "other",
    max_race_p == pred.whi ~ "white",
    TRUE ~ NA_character_
  ))
examiner_race

## # A tibble: 3,806 × 8
##   surname    pred.whi pred.bla pred.his pred.asi pred.oth max_race_p race
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <chr>
## 1 HOWARD      0.597    0.295    0.0275   0.00690   0.0741    0.597    white
## 2 YILDIRIM    0.807    0.0273   0.0694   0.0165    0.0798    0.807
```

```

white
## 3 HAMILTON      0.656  0.239  0.0286  0.00750  0.0692  0.656
white
## 4 MOSHER        0.915  0.00425 0.0291  0.00917  0.0427  0.915
white
## 5 BARR          0.784  0.120  0.0268  0.00830  0.0615  0.784
white
## 6 GRAY          0.640  0.252  0.0281  0.00748  0.0724  0.640
white
## 7 MCMILLIAN     0.322  0.554  0.0212  0.00340  0.0995  0.554
black
## 8 FORD          0.576  0.320  0.0275  0.00621  0.0697  0.576
white
## 9 STRZELECKA    0.472  0.171  0.220  0.0825  0.0543  0.472
white
## 10 KIM           0.0169 0.00282 0.00546 0.943  0.0319  0.943
Asian
## # ... with 3,796 more rows

```

Let's join the data back to the applications table.

```

# removing extra columns
examiner_race <- examiner_race %>%
  select(surname,race)
applications <- applications %>%
  left_join(examiner_race, by = c("examiner_name_last" = "surname"))
rm(examiner_race)
rm(examiner_surnames)
gc()

##           used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  4806713 256.8   8229292 439.5   8229292 439.5
## Vcells 54213238 413.7  111746554 852.6  92935921 709.1

```

Examiner's tenure

To figure out the timespan for which we observe each examiner in the applications data, let's find the first and the last observed date for each examiner. We'll first get examiner IDs and application dates in a separate table, for ease of manipulation. We'll keep examiner ID (the field `examiner_id`), and earliest and latest dates for each application (`filing_date` and `appl_status_date` respectively). We'll use functions in package `lubridate` to work with date and time values.

```

library(lubridate) # to work with dates
examiner_dates <- applications %>%
  select(examiner_id, filing_date, appl_status_date)
examiner_dates

## # A tibble: 2,018,477 × 3
##   examiner_id filing_date appl_status_date
##         <dbl> <date>      <chr>

```

```
## 1      96082 2000-01-26 30jan2003 00:00:00
## 2      87678 2000-10-11 27sep2010 00:00:00
## 3      63213 2000-05-17 30mar2009 00:00:00
## 4      73788 2001-07-20 07sep2009 00:00:00
## 5      77294 2000-04-10 19apr2001 00:00:00
## 6      68606 2000-04-28 16jul2001 00:00:00
## 7      89557 2004-01-26 15may2017 00:00:00
## 8      97543 2000-06-23 03apr2002 00:00:00
## 9      98714 2000-02-04 27nov2002 00:00:00
## 10     65530 2002-02-20 23mar2009 00:00:00
## # ... with 2,018,467 more rows
```

The dates look inconsistent in terms of formatting. Let's make them consistent. We'll create new variables `start_date` and `end_date`.

```
examiner_dates <- examiner_dates %>%
  mutate(start_date = ymd(filing_date), end_date =
as_date(dmy_hms(appl_status_date)))
```

Let's now identify the earliest and the latest date for each examiner and calculate the difference in days, which is their tenure in the organization.

```
examiner_dates <- examiner_dates %>%
  group_by(examiner_id) %>%
  summarise(
    earliest_date = min(start_date, na.rm = TRUE),
    latest_date = max(end_date, na.rm = TRUE),
    tenure_days = interval(earliest_date, latest_date) %/% days(1)
  ) %>%
  filter(year(latest_date)<2018)
examiner_dates
```

```
## # A tibble: 5,625 × 4
##   examiner_id earliest_date latest_date tenure_days
##   <dbl> <date> <date> <dbl>
## 1      59012 2004-07-28 2015-07-24 4013
## 2      59025 2009-10-26 2017-05-18 2761
## 3      59030 2005-12-12 2017-05-22 4179
## 4      59040 2007-09-11 2017-05-23 3542
## 5      59052 2001-08-21 2007-02-28 2017
## 6      59054 2000-11-10 2016-12-23 5887
## 7      59055 2004-11-02 2007-12-26 1149
## 8      59056 2000-03-24 2017-05-22 6268
## 9      59074 2000-01-31 2017-03-17 6255
## 10     59081 2011-04-21 2017-05-19 2220
## # ... with 5,615 more rows
```

Joining back to the applications data.

```
applications <- applications %>%
  left_join(examiner_dates, by = "examiner_id")
```

```
rm(examiner_dates)
gc()

##           used (Mb) gc trigger (Mb) max used (Mb)
## Ncells  4813327 257.1  14938879 797.9  14938879 797.9
## Vcells 64477940 492.0  134175864 1023.7 133649394 1019.7
```

Examiner's workgroup

Let's create workgroup for each examiner in the applications data

```
applications$workgroups <- substr(applications$examiner_art_unit, 1, 3)
```

Count the number of examiners in each workgroup

```
workgroups_count = applications %>%
  distinct(examiner_id, workgroups) %>%
  group_by(workgroups) %>%
  summarize(count = n()) %>%
  arrange(desc(count))
```

We will choose the top 2 workgroups with the largest number of examiner, which is 179 and 216

```
# Create a new table for applications with `workgroups` equal to 179
workgroup_179 <- applications %>%
  filter( workgroups == "179")

workgroup_179 <- distinct(workgroup_179, examiner_id, .keep_all = TRUE)


# Create a new table for applications with `workgroups` equal to 216
workgroup_216 <- applications %>%
  filter( workgroups == "216")

workgroup_216 <- distinct(workgroup_216, examiner_id, .keep_all = TRUE)
```

We will examine the demography of both workgroup

Workgroup 179:

```
ggplot(workgroup_179, aes(x = race, fill = gender)) +
  geom_bar(position = "dodge") +
  xlab("Race") +
  ylab("Count") +
  ggtitle("Distribution of Gender and Race in Workgroup 179")
```


```
table(workgroup_179$gender)
```

```
##
## female   male
##    308    523
```

```
table(workgroup_179$race)
```

```
##
##   Asian   black Hispanic   other   white
##    200     24      23      1    666
```

```
ggplot(workgroup_216, aes(x = race, fill = gender)) +
  geom_bar(position = "dodge") +
  xlab("Race") +
  ylab("Count") +
  ggtitle("Distribution of Gender and Race in Workgroup 216")
```



```
table(workgroup_216$gender)
```

```
##
## female   male
##      80    222
```

```
table(workgroup_216$race)
```

```
##
##   Asian   black Hispanic   white
##    129     18      12     200
```

In terms of gender, both groups have more male than female. In addition, if we look at the composition of gender within a specific race, it is indicated that for both Asian and White, male examiners outnumber their female peers.

In terms of race, Asian and White are the two dominating race in both workgroup. For 179, the number of white examiner is approximately 3.3 more than that of Asian examiner. Similarly, in 216, white examiners outnumber their Asian peers by 71 examiners. For both group, Black and Hispanic are underrepresented.

Network between examiner

We need to create a node dataframe with the examiner ID

```
#Create the nodes dataframe
nodes <- applications %>%
```

```
distinct(examiner_id) %>%
select(examiner_id)
```

nodes

```
## # A tibble: 5,649 × 1
##   examiner_id
##   <dbl>
## 1      96082
## 2      87678
## 3      63213
## 4      73788
## 5      77294
## 6      68606
## 7      89557
## 8      97543
## 9      98714
## 10     65530
## # ... with 5,639 more rows
```

We then proceed to create a new edges dataframe by cleaning the original edges dataframe

```
#Subset examiner_id and workgroup columns from applications
examiner_workgroup <- subset(applications, select = c("examiner_id",
"workgroups"))

# Look up the workgroup corresponding to each ego_examiner_id in the
examiner_workgroup dataframe
edges$ego_examiner_workgroup <-
examiner_workgroup$workgroups[match(edges$ego_examiner_id,
examiner_workgroup$examiner_id)]

# Look up the workgroup corresponding to each alter_examiner_id in the
examiner_workgroup dataframe
edges$alter_examiner_workgroup <-
examiner_workgroup$workgroups[match(edges$alter_examiner_id,
examiner_workgroup$examiner_id)]

# Filter out observations where either ego_examiner_workgroup or
alter_examiner_workgroup is NA
edges_filtered <- edges[complete.cases(edges[, c("ego_examiner_id",
"alter_examiner_id", "ego_examiner_workgroup", "alter_examiner_workgroup"))],
]

#Create the edges dataframe

edges_1 <- edges_filtered %>%
  distinct(ego_examiner_id, alter_examiner_id) %>%
  select(ego_examiner_id, alter_examiner_id)
```

```
edges_1
## # A tibble: 6,385 × 2
##   ego_examiner_id alter_examiner_id
##           <dbl>           <dbl>
## 1           84356           66266
## 2           84356           63519
## 3           84356           98531
## 4           92953           93865
## 5           92953           91818
## 6           72253           61519
## 7           72253           72253
## 8           72253           67515
## 9           67078           75772
## 10          67078           97328
## # ... with 6,375 more rows
```

Let's create a network graph between the examiners

```
library(igraph)

##
## Attaching package: 'igraph'

## The following objects are masked from 'package:lubridate':
##
##   %--%, union

## The following objects are masked from 'package:dplyr':
##
##   as_data_frame, groups, union

## The following objects are masked from 'package:purrr':
##
##   compose, simplify

## The following object is masked from 'package:tidyr':
##
##   crossing

## The following object is masked from 'package:tibble':
##
##   as_data_frame

## The following objects are masked from 'package:stats':
##
##   decompose, spectrum

## The following object is masked from 'package:base':
##
##   union
```

```
# Create network graph
```

```
g <- graph_from_data_frame(d = edges_1, directed = FALSE, vertices = nodes)
```

```
## Warning in graph_from_data_frame(d = edges_1, directed = FALSE, vertices =  
## nodes): In `vertices[,1]' `NA' elements were replaced with string "NA"
```

For the purpose of this project, I will measure degree centrality. With the large number of employees in the organization, the person with the most connections will be influential as their connection could help them in their projects. Looking from a project perspective, knowing a lot of people will give you an advantage.

Let's calculate the degree centrality for all examiner_id

```
# Calculate the degree centrality of each node (examiner ID)
```

```
centrality <- degree(g, mode = "all", normalized = FALSE)
```

```
# Combine the centrality scores and node IDs into a table
```

```
centrality_table <- data.frame(examiner_id = V(g)$name, centrality =  
as.vector(centrality))
```

```
centrality_table$examiner_id <- as.numeric(centrality_table$examiner_id)
```

```
## Warning: NAs introduced by coercion
```

Let's calculate the degree centrality for each examiner in group 179

```
workgroup_179 <- left_join(workgroup_179, centrality_table, by =  
"examiner_id")
```

```
workgroup_179 <- arrange(workgroup_179, desc(centrality))
```

```
workgroup_179
```

```
## # A tibble: 914 × 23
```

```
##   applicat...1 filing_d...2 exami...3 exami...4 exami...5 exami...6 exami...7 uspc...8
```

```
uspc...9
```

```
##   <chr>      <date>      <chr>    <chr>    <chr>      <dbl>    <dbl> <chr>  
<chr>
```

```
## 1 09902475   2001-07-09 DAVIS    ROBERT   B          63176    1791 425  
195000
```

```
## 2 09717374   2000-11-22 HENDRI... STUART   L          67698    1793 502  
417000
```

```
## 3 10531525   2005-11-02 INYARD   APRIL    C          79847    1794 428  
221000
```

```
## 4 09727516   2000-12-04 ALANKO   ANITA    KAREN     71119    1792 216  
083000
```

```
## 5 09849065   2001-05-04 PADGETT MARIAN... L          75341    1792 427  
569000
```

```
## 6 09731945   2000-12-07 DANIELS  MATTHEW  J          59771    1791 264  
257000
```

```
## 7 10399396   2003-04-21 JOY      DAVID    J          80730    1794 428  
292100
```

```
## 8 10362305   2003-05-21 PARVINI  PEGAH    <NA>      66436    1793 106  
487000
```

```
## 9 10402745 2003-03-28 DAHIME... MAHMOUD <NA> 93432 1792 438
710000
## 10 10523571 2006-02-27 BELL WILLIAM P 76154 1791 264
243000
## # ... with 904 more rows, 14 more variables: patent_number <chr>,
## # patent_issue_date <date>, abandon_date <date>, disposal_type <chr>,
## # appl_status_code <dbl>, appl_status_date <chr>, tc <dbl>, gender
<chr>,
## # race <chr>, earliest_date <date>, latest_date <date>, tenure_days
<dbl>,
## # workgroups <chr>, centrality <dbl>, and abbreviated variable names
## # 1application_number, 2filing_date, 3examiner_name_last,
## # 4examiner_name_first, 5examiner_name_middle, 6examiner_id, ...
```

Let's calculate the degree centrality for each examiner in group 216

```
workgroup_216 <- left_join(workgroup_216, centrality_table, by =
"examiner_id")
workgroup_216 <- arrange(workgroup_216, desc(centrality))
workgroup_216

## # A tibble: 359 × 23
## applicat...1 filing_d...2 exami...3 exami...4 exami...5 exami...6 exami...7 uspc_...8
uspc_...9
## <chr> <date> <chr> <chr> <chr> <dbl> <dbl> <chr>
<chr>
## 1 09564248 2000-05-04 PHAM KHANH B 78326 2166 707
204000
## 2 09580327 2000-05-26 WASSUM LUKE S 99902 2167 707
010000
## 3 10481715 2004-06-01 WONG JOSEPH D 93865 2166 707
104100
## 4 09602441 2000-06-23 ALAM HOSAIN T 92060 2166 707
101000
## 5 09507064 2000-02-18 ABEL J... NEVEEN <NA> 83222 2165 707
100000
## 6 11071937 2005-03-04 LIAO JASON G 70672 2169 707
002000
## 7 09782596 2001-02-12 STEVENS ROBERT <NA> 71230 2162 715
517000
## 8 09479669 2000-01-10 COBY FRANTZ <NA> 60130 2161 707
007000
## 9 09656484 2000-09-07 LU KUEN S 83313 2167 707
010000
## 10 09505165 2000-02-16 ALI MOHAMM... <NA> 67404 2167 707
001000
## # ... with 349 more rows, 14 more variables: patent_number <chr>,
## # patent_issue_date <date>, abandon_date <date>, disposal_type <chr>,
## # appl_status_code <dbl>, appl_status_date <chr>, tc <dbl>, gender
<chr>,
```

```
## #   race <chr>, earliest_date <date>, latest_date <date>, tenure_days  
<dbl>,  
## #   workgroups <chr>, centrality <dbl>, and abbreviated variable names  
## #   1application_number, 2filing_date, 3examiner_name_last,  
## #   4examiner_name_first, 5examiner_name_middle, 6examiner_id, ...
```

In group 179, the top 10 examiners with the highest centrality score are all white, with 4 females and 6 males. This could be explained by the fact that white is the dominating race in the group, and subconsciously people would make connection with people with similar backgrounds, in this case, of white race. On the other hand, in group 216, the top 10 examiners with the highest centrality score is mostly minorities (6 Asians, 1 black). This is interesting as Asian is not the major race in the group. This could be due to the difference in the topic that each workgroup deal with, and the people with the highest centrality score are those that are good at these topics.